

# OpenCL Ray Tracer

## Ray Tracing on the GPU

Xinyu Cheng, Jake Crouch, Spencer Murphy, and Dylan Sturgeon



# Roadmap

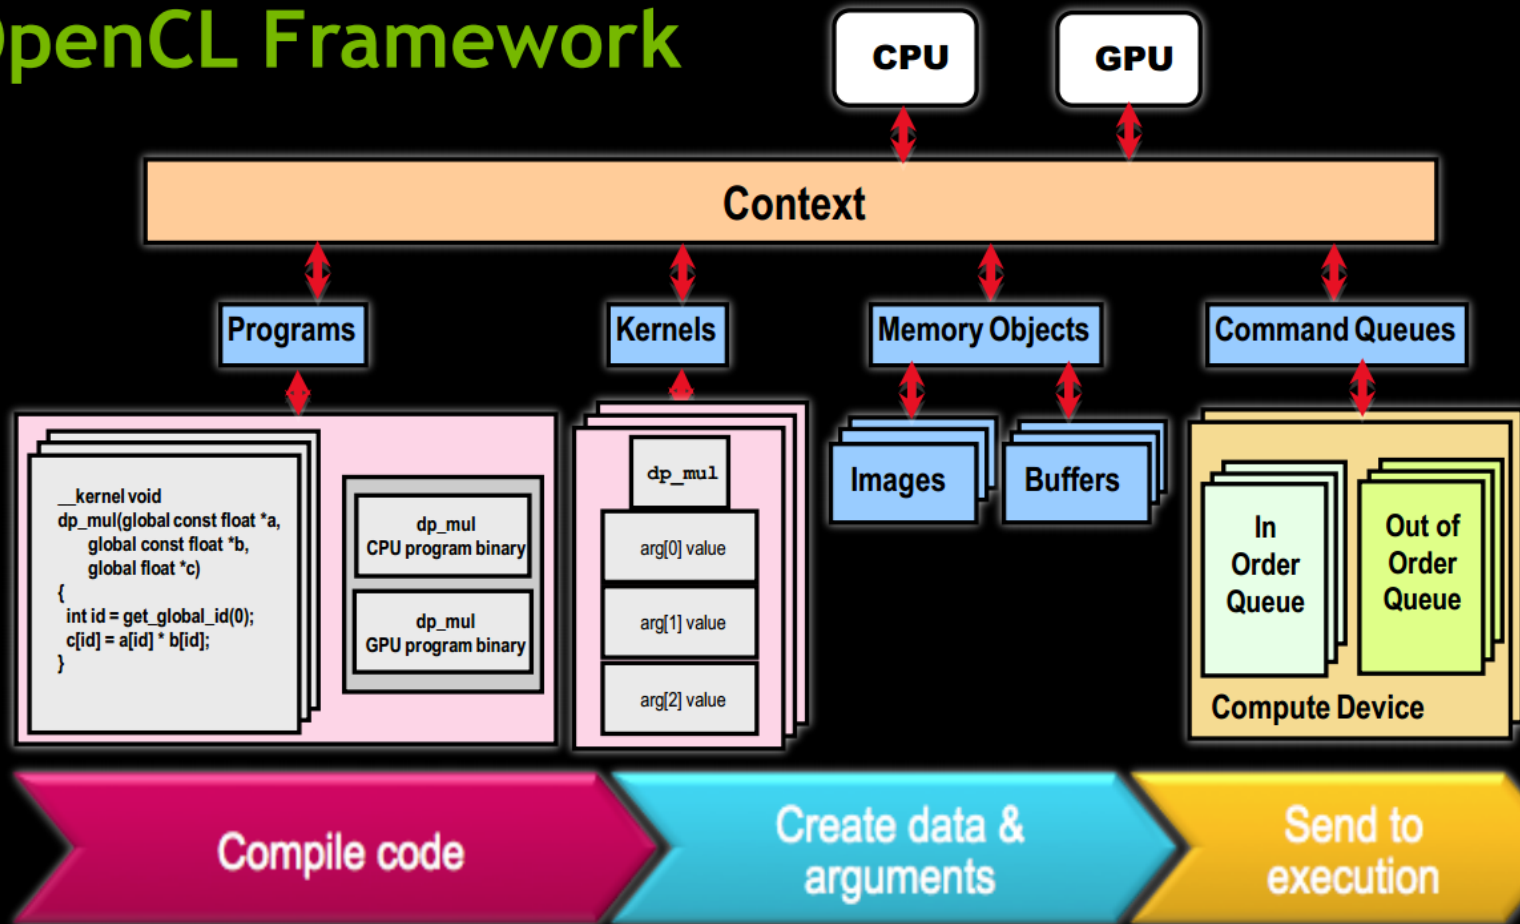
- OpenCL Research
- GPU Ray Tracing
- Results
- Future Work

# OpenCL - OpenWhat?



- Open Computing Language
  - Standard for Parallel Programming
  - Maintained by Khronos Group
  - Open and royalty free
- Implemented by most modern hardware
  - All Intel/AMD CPUs
  - All NVidia/AMD GPUs

# OpenCL Framework



# OpenCL - C Code

Derivative of ISO C99

Additional Features

- Vector Types (float3, float4, etc.)
- Address Space qualifiers
- Synchronization
- Built-in functions

# OpenCL - Code Restrictions

Pointers to functions

Pointers as an argument

Recursion

Dynamic memory

# OpenCL - Kernels

Build a kernel from file

Create memory objects

Pass arguments and execute kernel

Read back memory objects

# OpenCL - Command Queues

## Commands Enqueued

- Execute Asynchronously (optionally block)

- In-order or Out-of-order (in-order by default)

## Command Queue

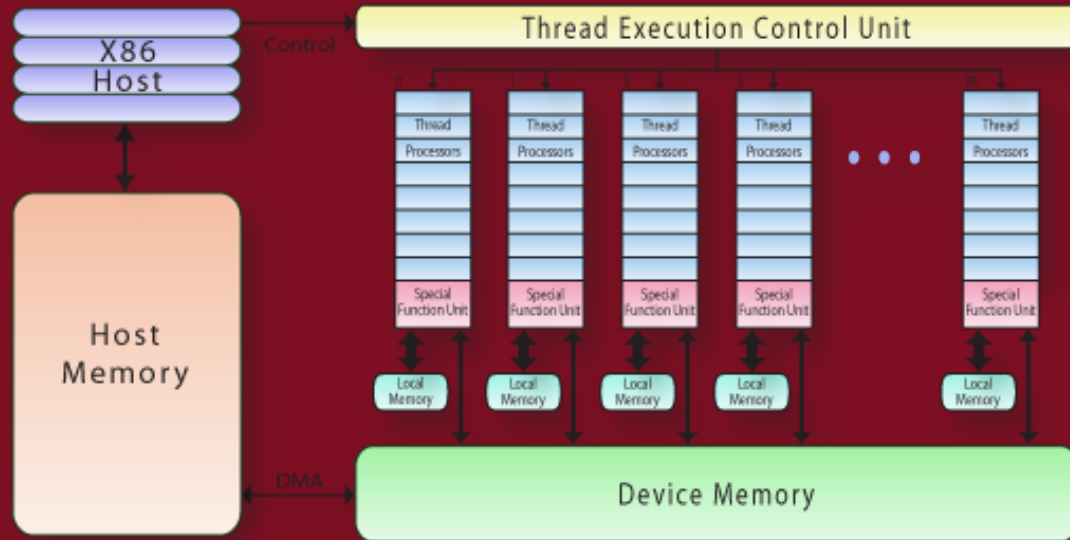
- Commands scheduled to execute on a device



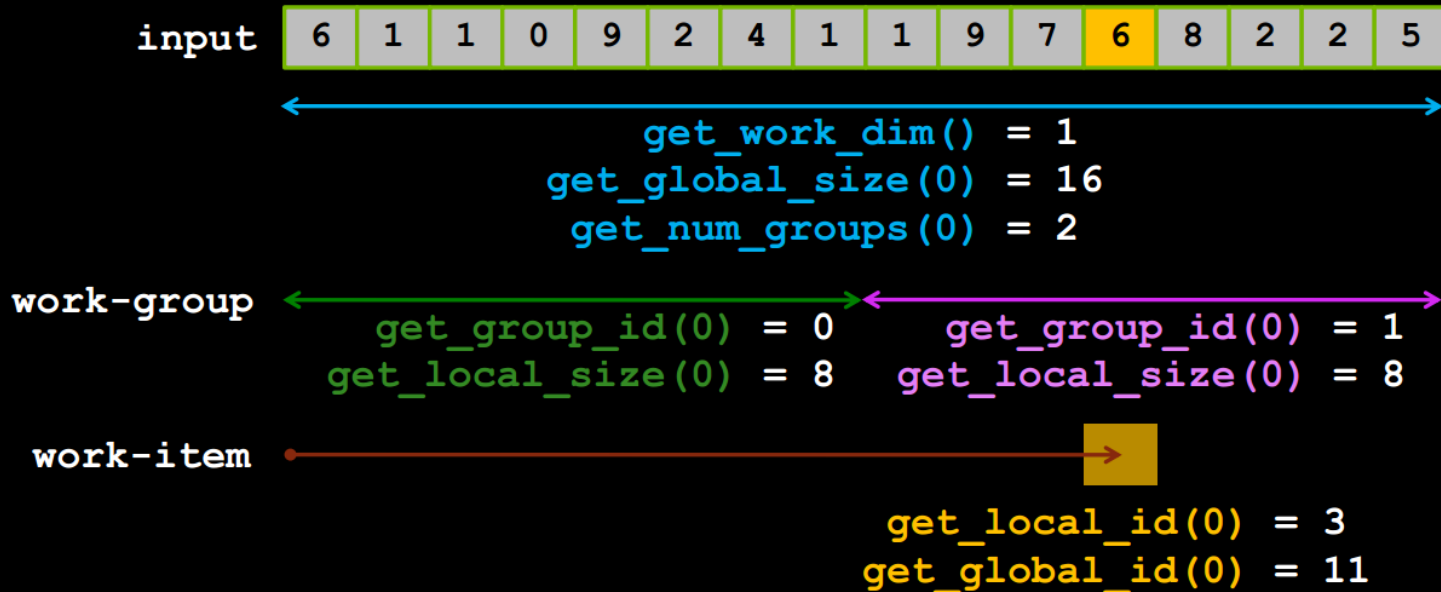
# OpenCL - Kernel Execution

Work-items execute kernels in lock step

Work-group - collection of work-items



# OpenCL - Work Group Example



# OpenCL - Errors

Any operation can (and probably will) fail!  
Operations set or return an error code  
Error codes map to *\*useful\** descriptions


Lesson: *Always* check for errors.

```
#define CL_INVALID_VALUE -30
```

Thanks Khronos...

# GPU Ray Tracing - Work in CPU & GPU

- Works in CPU
  - Load objects and materials.
  - Build camera, scene and tree.
- Works in GPU
  - Run intersections.
  - Calculate color.



# GPU Ray Tracing - Data Passing

## Basic Idea for Data Passing

In CPU, we create buffer for input and output of GPU functions. The buffer data can be used in the device functions.

```
Camera camera = scene.camera;
cl_mem cameraBuffer = clKernel.createBuffer(CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR, (1) * sizeof(Camera), (void *)&camera, &err);
clKernel.checkErr(err, "creating camera buffer");
```

In GPU, we create the corresponding functions to receive input buffers. Then do the calculation and put results into output buffers.

```
__kernel void calculateHitPoints(__global Node* nodes, int nodeCount, __global Triangle* triangles,
    int trianglesSize, __global Camera* camera, int width, int height, __global HitPoint* hitPoints)
{
```

```
err = clKernel.readBuffer(hitPoints, CL_TRUE, 0, width*height * sizeof(HitPoint), (void*)&outHits[0], 0, NULL, NULL);
clKernel.checkErr(err, "reading buffer");
```

# GPU Ray Tracing - Multi Kernel

Two kernels:

Intersection Kernel and Color Kernel.

Intersection Kernel:

```
ClKernel clKernel("HitPointCalculator.cl", CL_DEVICE_TYPE_GPU, "calculateHitPoints");
```

```
__kernel void calculateHitPoints(__global Node* nodes, int nodeCount, __global Triangle* triangles,  
    int trianglesSize, __global Camera* camera, int width, int height, __global HitPoint* hitPoints)
```

For intersection, we pass tree nodes, triangles, camera and scene resolution to the kernel function. The output will be put into hitPoints, which is a list of hit points we get from all intersections.

Color Kernel:

```
ClKernel clKernel("MaterialShaderKernel.cl", CL_DEVICE_TYPE_GPU, "calculateMaterialColors");
```


```
__kernel void calculateMaterialColors(__global HitPoint* hitPoints, __global Material* materials, __global Light* lights, int lightSize, __global Triangle* triangles,  
    int triangleSize, __global Node* nodeLists, int nodeCount, __global float* colors, float cameraOriginX, float cameraOriginY, float cameraOriginZ)
```

For color calculation, we pass the hitPoints, materials, lights, triangles and nodelists to the GPU. The output is the colors.

# GPU Ray Tracing - No Recursion

Recursion does not work in GPU  
Create a Stack for tree traversal  
Use while loop for reflections

We made fractals too!  
But only in 3D...




# GPU Ray Tracing - Flat Trees

Can't pass pointers into a kernel

Flat Tree - Tree as list of nodes

Traverse using list indices






# Results


It's fast

But could be faster...

Approximately 10x fps of  
original ray tracer




# Results - Cornell Box




OpenCL - GPU					C++ (CPU)				
Resolution	256	512	1024	2048	Resolution	256	512	1024	2048
Run Time (ms)	6	13	45	183	Run Time (ms)	121	471	1737	7240
FPS	166.7	76.9	22.2	5.5	FPS	8.36	2.12	0.58	0.14

# Results - Blue Sphere



OpenCL - GPU					C++ (CPU)				
Resolution	256	512	1024	2048	Resolution	256	512	1024	2048
Run Time (ms)	10	29	99	333	Run Time (ms)	97	375	1459	5521
FPS	100	34.5	10.1	3.00	FPS	10.3	2.67	0.69	0.18

# Results - Dragon



OpenCL - GPU					C++ (CPU)				
Resolution	256	512	1024	2048	Resolution	256	512	1024	2048
Run Time (ms)	34	118	468	1580	Run Time (ms)	185	648	2311	8323
FPS	29.4	8.47	2.14	0.63	FPS	5.41	1.54	0.432	0.120



# Future Work

## Optimizations

- Our tree used spatial splitting - there are better algorithms
- Improve memory reference locality
  - Use less global memory
  - Kernels are faster when referencing contiguous chunks

# Future Work (cont.)

## Animations

- Providing movement to the scene.
- Can move the camera to create a 360 degree view of the static scene but would like to add additional translations.

## Textures

- Allow for drawing images instead of only grabbing visual information from the material

# Questions?

## Sample Questions

- What is this Ray Tracer thing?
- Does the real world application of Ray Tracing use the GPU or the CPU?
- Where's the Millennium Falcon?